

Reinforcement learning (I)

AI 2025/2026

Content

Introduction

Markov decision process

Value iteration

Policy iteration

Reinforcement learning

The agent must learn a behavior, without having an instructor


- ▶ The agent has a task to perform
- ▶ Performs a series of actions
- ▶ Receives *feedback* from the environment: how well he acted to accomplish the task.

The agent receives a positive reward if the task is performed well, otherwise a negative reward.


This learning method is called **reinforcement learning**.

The interaction model

- ▶ The agent performs **actions**
- ▶ The environment presents the agent situations called **states** and gives **rewards**




Trajectory / episode $(s_0, a_0, r_1, s_1, a_1, r_2, \dots)$



Reinforcement learning

- ▶ The goal: to make the agent to act in order to maximize his rewards
- ▶ The agent must identify the **sequence of actions** leading to the completion of the task
 - ▶ Data acquisition and supervised learning on data
Training data: (S, A, R) State, Action, Reward

Applications: TD-Gammon



A neural network trained with *TD-learning*

- ▶ input: configuration
- ▶ output: an estimation of the value for that configuration

Learn from simulations (play games against himself)

Examples: robotics

AIBO: learn to walk '04 (using *Policy gradient*)



Fig. 5. The training environment for our experiments. Each Aibo times itself as it moves back and forth between a pair of landmarks (A and A', B and B', or C and C').

[https://www.cs.utexas.edu/users/
AustinVilla/?p=research/learned_walk](https://www.cs.utexas.edu/users/AustinVilla/?p=research/learned_walk)

Minitaur quadrupedal robot: learn to walk (using *actor-critic deep RL*)




Fig. 9: We trained the Minitaur robot to walk on flat terrain (first row) in about two hours. At test time, we introduced obstacles, including a slope, wooden blocks, and steps, which were not present at training time, and the learned policy was able to generalize to the unseen situations without difficulty (other rows).

[https://sites.google.com/view/
minitaur-locomotion/](https://sites.google.com/view/minitaur-locomotion/)

Deep Reinforcement Learning

AlphaGo (Google DeepMind, '15): the program learned to play Atari 2600 games by directly watching only the display and the score.

- ▶ 2016: won against a 9 dan professional GO player (prize of 1 000 000\$)
- ▶ 2017: won against Ke Jie, the best GO player in the world

AlphaGo Zero learned to play **without input from human games** (just based on the rules of the game) and beat AlphaGo with 100-0 (Oct. '17)

AlphaZero beat AlphaGo Zero with 60-40, and after just 8 hours of training outperformed all existing **GO** and **CHESS** programs (Dec. '17)

Content



Introduction

Markov decision process

Value iteration

Policy iteration


Sequential decisions



- ▶ Deterministic environment
 - ▶ (up, up, right, right, right)
- ▶ Stochastic environment
 - ▶ **Transitions model** $P(s'|s, a)$: probability of reaching state s' if action a is done in state s
 - ▶ The action obtains the intended effect with probability 0.8
 - ▶ The agent receives a **reward**: -0.04 for nonterminal states; +/-1 for terminal states

Deterministic vs. stochastic

Stochastic: for the sequence (up, up, right, right, right), the probability of obtaining the final state is $0.8^5 = 0.33$



Markov assumption

- ▶ The current state s_t depends on a finite history of previous states
- ▶ First order Markov process: the current state s_t depends only on the previous state s_{t-1}

$$P(s_t | s_{t-1}, \dots, s_0) = P(s_t | s_{t-1})$$

Markov Decision Process (MDP)


Markov Decision Process: a sequential decision problem for a stochastic environment with a Markov transition model and additive rewards

- ▶ States $s \in S$ (initial state s_0), actions $a \in A$
- ▶ Transitions model $P(s'|s, a)$
- ▶ Reward function $R(s)$

What does a solution look like? Must specify what the agent should do in each state (**policy π**).


MDPs are non-deterministic search problems.

- ▶ Search algorithms. Expectimax alg: when the opponent does not play optimally.

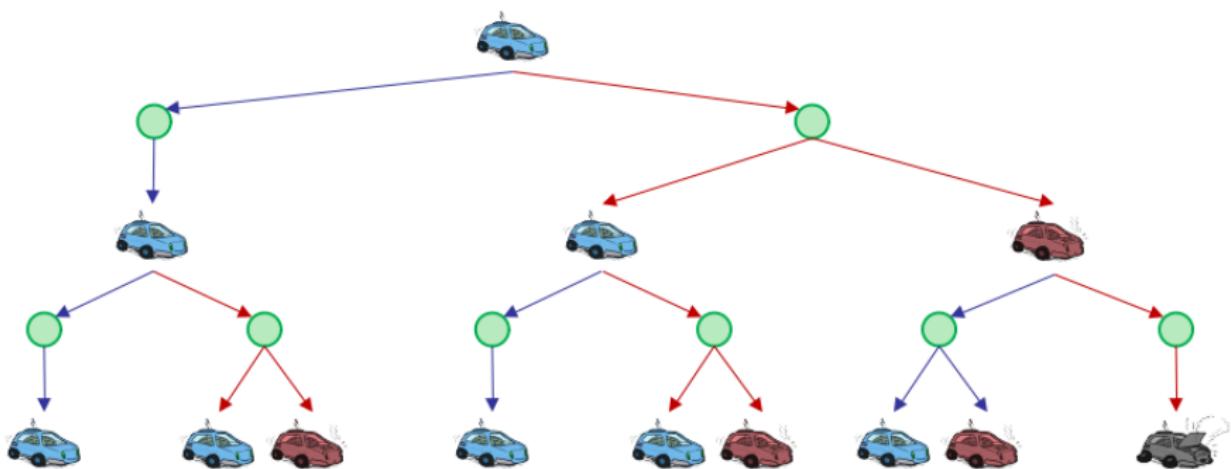


Example

- ▶ A robot car wants to travel far and quickly.
- ▶ States: Cool, Warm, Overheated; Actions: Slow, Fast; Going faster gets double reward.




Example: search tree



MDP search trees

- ▶ Each state projects a search tree



s state, (s, a, s') transition, $P(s'|s, a)$, $R(s, a, s')$

- ▶ In search problems, the goal is to identify an optimal *sequence*.

In MDP, the goal is to identify an *optimal policy* π^* (strategy)


- ▶ $\pi : S \rightarrow A$

$\pi(s)$ is the action recommended in state s

- ▶ **Utility:** the sum of rewards for a **sequence** of states.
The reward is the immediate, short-term gain; utility is the total, long-term gain.
- ▶ **Quality of a policy:** **expected utility** of possible sequences of states.
Stochastic environment: we can have a different sequence of states when executing the same policy from the initial state.

Optimal policy π^* maximizes the expected utility.

Example: the optimal policy and the state values




3	0.812	0.868	0.912	+1
2	0.762		0.660	-1
1	0.705	0.655	0.611	0.388

Utilities

Finite horizon

- ▶ $U_h([s_0, s_1, \dots, s_{N+k}]) = U_h([s_0, s_1, \dots, s_N]), \forall k > 0$
After a fixed time N , nothing matters.
- ▶ The optimal policy is not stationary: the optimal action for a given state may change over time

Example:



- ▶ $N = 3 \rightarrow$ have to risk (up)
- ▶ $N = 100 \rightarrow$ can choose the safer solution (left)

Utilities

Infinite horizon

- ▶ There is no fixed deadline
- ▶ The optimal policy is stationary
- ▶ a. Additive rewards

$$U_h([s_0, s_1, s_2, \dots]) = R(s_0) + R(s_1) + R(s_2) + \dots$$

- ▶ b. *Discounted* rewards

$$U_h([s_0, s_1, s_2, \dots]) = R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots$$

$\gamma \in [0, 1]$ *discount factor* indicates that future rewards matter less than immediate ones

Discounted rewards

$$U_h([s_0, s_1, s_2, \dots]) = R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots$$

- ▶ Prefer current rewards, not later ones
- ▶ Values of rewards decay exponentially

Example 1: $\gamma = 0.5$

$$R(s_0) = 1, R(s_1) = 2, \dots$$

$$U([1, 2, 3]) = 1 * 1 + 0.5 * 2 + 0.25 * 3$$

$$U([1, 2, 3]) < U([3, 2, 1])$$

Example 2: Quiz

Infinite horizon - evaluation

We must ensure that the utility of a possibly infinite sequence is finite.

- ▶ **1st Approach.** If the rewards are bounded and $\gamma < 1$ then:

$$U_h([s_0, s_1, s_2, \dots]) = \sum_{t=0}^{\infty} \gamma^t R(s_t) \leq \sum_{t=0}^{\infty} \gamma^t R_{max} = R_{max}/(1 - \gamma)$$

- ▶ **2nd Approach.** If the environment contains terminal states and it's guaranteed that the agent will reach one of them (we have a *proper policy*), we can use $\gamma = 1$

Expected utility

- ▶ Each policy generates multiple sequences of states, due to the uncertainty of transitions $P(s'|s, a)$
- ▶ Let S_t be a **random variable**: the state the agent reaches at time t executing policy π ; $S_0 = s$.

Expected utility obtained by executing policy π starting in state s :

$$U^\pi(s) = E \left[\sum_{t=0}^{\infty} \gamma^t R(S_t) \right]$$

(= the expected value of the sum of all discounted rewards obtained for all possible sequences of states)

Evaluation of a policy

- ▶ Optimal policy

$$\pi_s^* = \operatorname{argmax}_\pi U^\pi(s)$$

- ▶ $U^{\pi^*}(s)$ the true utility of a state: the expected value of the sum of the discounted rewards if the agent executes an optimal policy

Example: Consider $\gamma = 1$ and $R(s) = -0.04$.

3	0.812	0.868	0.912	+1
2	0.762		0.660	-1
1	0.705	0.655	0.611	0.388

Near the final state, the utilities are higher because fewer steps with negative reward are required to reach that state.

Maximum Expected Utility

Maximum Expected Utility: choose the action that maximizes the expected utility of the subsequent state

$$\pi^*(s) = \operatorname{argmax}_{a \in A(s)} \sum_{s'} P(s'|s, a) U(s')$$

Bellman equation

$$U(s) = R(s) + \gamma \max_{a \in A(s)} \sum_{s'} P(s'|s, a) U(s')$$

Bellman equation (1957): the utility of a state is the immediate reward for that state, $R(s)$, plus the maximum expected utility of the next state.

Example

The utility of state (1,1):

$$U(1,1) = -0.04 + \gamma \max[0.8U(1,2) + 0.1U(2,1) + 0.1U(1,1), \quad (Up)$$
$$\qquad\qquad\qquad 0.9U(1,1) + 0.1U(1,2), \quad (Left)$$
$$\qquad\qquad\qquad 0.9U(1,1) + 0.1U(2,1), \quad (Down)$$
$$\qquad\qquad\qquad 0.8U(2,1) + 0.1U(1,2) + 0.1U(1,1)] \quad (Right)$$

The best action: *Up*.

Solving a Markov decision process

- ▶ n possible states
- ▶ n Bellman equations, one for each state
- ▶ n equations with n unknowns: $U(s)$
- ▶ Cannot be solved as a system of linear equations because of the *max* function

Content

Introduction

Markov decision process

Value iteration

Policy iteration

I. Value iteration

Computes the utility of each state and identifies the optimal action in each state.

The algorithm for computing the optimal policy:

- ▶ Initializes utilities with arbitrary values
- ▶ Updates the utility of each state from the utilities of its neighbors

$$U_{i+1}(s) = R(s) + \gamma \max_a \sum_{s'} P(s'|s, a) U_i(s')$$

- ▶ Repeat for each s simultaneously, until an equilibrium is reached

Value iteration: pseudocode

function VALUE-ITERATION(mdp, ϵ) **returns** a utility function

inputs: mdp , an MDP with states S , actions $A(s)$, transition model $P(s' | s, a)$, rewards $R(s)$, discount γ

ϵ , the maximum error allowed in the utility of any state

local variables: U, U' , vectors of utilities for states in S , initially zero

δ , the maximum change in the utility of any state in an iteration

repeat

$U \leftarrow U'; \delta \leftarrow 0$

for each state s **in** S **do**

$U'[s] \leftarrow R(s) + \gamma \max_{a \in A(s)} \sum_{s'} P(s' | s, a) U[s']$

if $|U'[s] - U[s]| > \delta$ **then** $\delta \leftarrow |U'[s] - U[s]|$

until $\delta < \epsilon(1 - \gamma)/\gamma$

return U

Value iteration

Value Iteration, for estimating $\pi \approx \pi_*$

Algorithm parameter: a small threshold $\theta > 0$ determining accuracy of estimation
Initialize $V(s)$, for all $s \in \mathcal{S}^+$, arbitrarily except that $V(\text{terminal}) = 0$

Loop:

```
|   Δ ← 0
|   Loop for each  $s \in \mathcal{S}$ :
|      $v \leftarrow V(s)$ 
|      $V(s) \leftarrow \max_a \sum_{s',r} p(s',r|s,a)[r + \gamma V(s')]$ 
|     Δ ← max(Δ, |v - V(s)|)
until Δ < θ
```

Output a deterministic policy, $\pi \approx \pi_*$, such that

$$\pi(s) = \arg \max_a \sum_{s',r} p(s',r|s,a)[r + \gamma V(s')]$$

From Sutton&Barto. Reinforcement Learning: an introduction

Utility function $U \leftrightarrow V$ value function

Value iteration

- ▶ Example: Racing
- ▶ Demo: <https://courses.grainger.illinois.edu/cs440/fa2018/lectures/mdp-value-demo.pdf>

Problems with Value iteration

- ▶ Slow: $O(S^2A)$ per iteration
- ▶ The “max” at each state rarely changes
- ▶ The policy often converges before the values

Content

Introduction

Markov decision process

Value iteration

Policy iteration

II. Policy iteration

- ▶ If we **fix the policy**, we have only one action per state
- ▶ The algorithm alternates the following steps:
 - ▶ 1. **Policy evaluation**: given a policy π_i , compute the utilities of the states based on the policy π_i : $U_i = U^{\pi_i}$
 - ▶ 2. **Policy improvement**: compute a new policy π_{i+1} , based on the utilities U_i

Repeat the steps until the policy converges

1. Policy evaluation

The action for each state is fixed by the policy; at iteration i , the policy π_i specifies the action $\pi_i(s)$ in state s .

The simplified Bellman equations

$$U_i(s) = R(s) + \gamma \sum_{s'} P(s'|s, \pi_i(s)) U_i(s')$$

- ▶ A system of n linear equations with n unknowns
- ▶ It can be solved *exactly* in $O(n^3)$ or *approximately*
- ▶ Apply *Value iteration*

$$U_{i+1}(s) = R(s) + \gamma \sum_{s'} P(s'|s, \pi_i(s)) U_i(s')$$

1. Policy evaluation

Example: $\pi_i(1, 1) = Up$, $\pi_i(1, 2) = Up$

Simplified Bellman equations:

$$U_i(1, 1) = -0.04 + 0.8U_i(1, 2) + 0.1U_i(1, 1) + 0.1U_i(2, 1)$$

$$U_i(1, 2) = -0.04 + 0.8U_i(1, 3) + 0.2U_i(1, 2)$$

2. Policy improvement

- ▶ The values $U(s)$ are known
- ▶ Computes for each s , the optimal action

$$a_i^*(s) = \max_a \sum_{s'} P(s'|s, a) U(s')$$

- ▶ If $a_i^*(s) \neq \pi_i(s)$, update the policy: $\pi_{i+1}(s) \leftarrow a_i^*(s)$
Only the "promising" parts of the search space can be updated.

Policy iteration: pseudocode

```
function POLICY-ITERATION(mdp) returns a policy
    inputs: mdp, an MDP with states  $S$ , actions  $A(s)$ , transition model  $P(s' | s, a)$ 
    local variables:  $U$ , a vector of utilities for states in  $S$ , initially zero
                     $\pi$ , a policy vector indexed by state, initially random

repeat
     $U \leftarrow \text{POLICY-EVALUATION}(\pi, U, mdp)$ 
    unchanged?  $\leftarrow$  true
    for each state  $s$  in  $S$  do
        if  $\max_{a \in A(s)} \sum_{s'} P(s' | s, a) U[s'] > \sum_{s'} P(s' | s, \pi[s]) U[s']$  then do
             $\pi[s] \leftarrow \operatorname{argmax}_{a \in A(s)} \sum_{s'} P(s' | s, a) U[s']$ 
            unchanged?  $\leftarrow$  false
    until unchanged?
return  $\pi$ 
```

Demo: https://cs.stanford.edu/people/karpathy/reinforcejs/gridworld_dp.html

Bibliography

- ▶ Russel&Norvig. Artificial Intelligence: A modern Approach. Ch. 17. Making Complex Decisions
- ▶ Sutton&Barto. Reinforcement Learning. An introduction Ch 3. Finite Markov Decision Processes, Ch 4. Dynamic Programming
<http://incompleteideas.net/book/RLbook2020.pdf>