

Reinforcement learning (II)

AI 2025/2026

Content

Introduction

Passive learning

Model-based learning: Adaptive dynamic programming

Model-free learning: Direct utility estimation

Model-free learning: Temporal-Difference learning

Active learning

Q-learning

Deep Q-learning

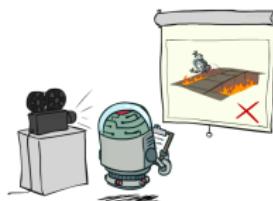
Conclusions

Reinforcement learning

- ▶ Markov decision process (MDP)
 - ▶ The set of states S , the set of actions A
 - ▶ The transition model $P(s'|s, a)$ is **known**
 - ▶ The reward function $R(s)$ is **known**
 - ▶ Computes an optimal policy
- ▶ Reinforcement learning
 - ▶ It is based on MDPs, but:
 - ▶ The transition model is **unknown**
 - ▶ The reward function is **unknown**
 - ▶ Learn an optimal policy

Types of reinforcement learning

- ▶ **Passive:** the agent learns the utility of being in certain states. The agent executes a **fixed** policy and **evaluates** it.



The agent has no control over his actions; applications: robotics.

- ▶ **Active:** the agent must learn what to do.



The agent must explore the environment and use the learned information. The agent updates its policy as it learns.

Types of reinforcement learning

- ▶ Model-based: learn the model of transitions and rewards and use it to discover the optimal policy
- ▶ Model-free: discover the optimal policy without learning the model

Content

Introduction

Passive learning

Model-based learning: Adaptive dynamic programming

Model-free learning: Direct utility estimation

Model-free learning: Temporal-Difference learning

Active learning

Q-learning

Deep Q-learning

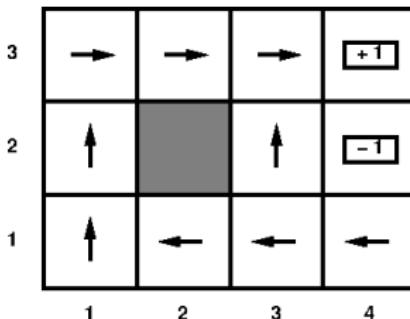
Conclusions

Passive learning

- ▶ The policy is **fixed**: always execute action $\pi(s)$ in state s
 - ▶ The goal: learn how good is the policy π
 - ▶ learn the utility $U^\pi(s)$ of each state
- how?
- execute the policy and learn from experience
- ▶ similar approach to step (1) policy evaluation from the *Policy iteration* algorithm; the difference: doesn't know the transition model $P(s'|s, a)$ and $R(s)$

Passive learning is a way of exploring the environment.

Passive learning



- ▶ The agent performs a series of trials (episodes)

$(1, 1)_{-.04} \rightsquigarrow (1, 2)_{-.04} \rightsquigarrow (1, 3)_{-.04} \rightsquigarrow (1, 2)_{-.04} \rightsquigarrow (1, 3)_{-.04} \rightsquigarrow (2, 3)_{-.04} \rightsquigarrow (3, 3)_{-.04} \rightsquigarrow (4, 3)_{+1}$

$(1, 1)_{-.04} \rightsquigarrow (1, 2)_{-.04} \rightsquigarrow (1, 3)_{-.04} \rightsquigarrow (2, 3)_{-.04} \rightsquigarrow (3, 3)_{-.04} \rightsquigarrow (3, 2)_{-.04} \rightsquigarrow (3, 3)_{-.04} \rightsquigarrow (4, 3)_{+1}$

$(1, 1)_{-.04} \rightsquigarrow (2, 1)_{-.04} \rightsquigarrow (3, 1)_{-.04} \rightsquigarrow (3, 2)_{-.04} \rightsquigarrow (4, 2)_{-1}$

- ▶ The policy is the same, but the environment is non-deterministic
- ▶ The goal is to learn the expected utility $U^\pi(s) = E [\sum_{t=0}^{\infty} \gamma^t R(S_t)]$

Content

Introduction

Passive learning

Model-based learning: Adaptive dynamic programming

Model-free learning: Direct utility estimation

Model-free learning: Temporal-Difference learning

Active learning

Q-learning

Deep Q-learning

Conclusions

Model-based learning: Adaptive dynamic programming (ADP)

1. Learn the transition model

Estimate $P(s'|s, \pi(s))$ and $R(s)$ from trials.

Use a table of probabilities: compute how often the result of an action occurs and estimate the transition probability.

Example: $(1, 1)_{-.04} \rightsquigarrow (1, 2)_{-.04} \rightsquigarrow (1, 3)_{-.04} \rightsquigarrow (1, 2)_{-.04} \rightsquigarrow (1, 3)_{-.04} \rightsquigarrow$
 $(2, 3)_{-.04} \rightsquigarrow (3, 3)_{-.04} \rightsquigarrow (4, 3)_{+1}$
 $(1, 1)_{-.04} \rightsquigarrow (1, 2)_{-.04} \rightsquigarrow (1, 3)_{-.04} \rightsquigarrow (2, 3)_{-.04} \rightsquigarrow (3, 3)_{-.04} \rightsquigarrow (3, 2)_{-.04} \rightsquigarrow$
 $(3, 3)_{-.04} \rightsquigarrow (4, 3)_{+1}$

$(1, 1)_{-.04} \rightsquigarrow (2, 1)_{-.04} \rightsquigarrow (3, 1)_{-.04} \rightsquigarrow (3, 2)_{-.04} \rightsquigarrow (4, 2)_{-1}$

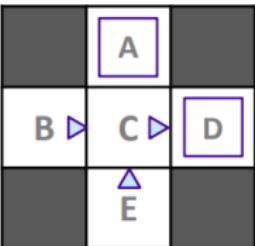
The action *Right* is executed 3 times in state $(1, 3)$ and in 2 cases the result is state $(2, 3)$ $\implies P((2, 3)|(1, 3), Right) = 2/3$

2. Solve MDP

Adaptive dynamic programming (ADP)

1. Learn the empiric model

Example:

Input Policy π	Observed (s, a, s', R) Transitions	Learned Model
	<p>Episode 1</p> <p>B, east, C, -1 C, east, D, -1 D, exit, x, +10</p> <p>Episode 2</p> <p>B, east, C, -1 C, east, D, -1 D, exit, x, +10</p>	$\hat{T}(s, a, s')$ T(B, east, C) = 1.00 T(C, east, D) = 0.75 T(C, east, A) = 0.25 ...
<i>Assume: $\gamma = 1$</i>	<p>Episode 3</p> <p>E, north, C, -1 C, east, D, -1 D, exit, x, +10</p> <p>Episode 4</p> <p>E, north, C, -1 C, east, A, -1 A, exit, x, -10</p>	$\hat{R}(s, a, s')$ R(B, east, C) = -1 R(C, east, D) = -1 R(D, exit, x) = +10 ...

Adaptive dynamic programming

2. Use dynamic programming to solve the MDP.

Enter the learned probabilities and rewards into the Bellman equations (fixed policy).

$$U^\pi(s) = R(s) + \gamma \sum_{s'} P(s'|s, \pi(s)) U^\pi(s')$$

Solve the system of linear equations with the unknowns $U^\pi(s)$.

ADP is inefficient if the state space is large

- ▶ system of linear equations of order n
- ▶ backgammon: 10^{20} equations with 10^{20} unknowns

Content

Introduction

Passive learning

Model-based learning: Adaptive dynamic programming

Model-free learning: Direct utility estimation

Model-free learning: Temporal-Difference learning

Active learning

Q-learning

Deep Q-learning

Conclusions

Model-free learning: Direct utility estimation

- ▶ The utility of a state is the total expected reward from that state forward (**reward-to-go**)

Example: $(1, 1)_{-.04} \rightsquigarrow (1, 2)_{-.04} \rightsquigarrow (1, 3)_{-.04} \rightsquigarrow (1, 2)_{-.04} \rightsquigarrow (1, 3)_{-.04} \rightsquigarrow$
 $(2, 3)_{-.04} \rightsquigarrow (3, 3)_{-.04} \rightsquigarrow (4, 3)_{+1}$
 $(1, 1)_{-.04} \rightsquigarrow (1, 2)_{-.04} \rightsquigarrow (1, 3)_{-.04} \rightsquigarrow (2, 3)_{-.04} \rightsquigarrow (3, 3)_{-.04} \rightsquigarrow$
 $(3, 2)_{-.04} \rightsquigarrow (3, 3)_{-.04} \rightsquigarrow (4, 3)_{+1}$
 $(1, 1)_{-.04} \rightsquigarrow (2, 1)_{-.04} \rightsquigarrow (3, 1)_{-.04} \rightsquigarrow (3, 2)_{-.04} \rightsquigarrow (4, 2)_{-1}$

The first attempt produces:

- ▶ in state $(1,1)$, total reward $0.72 (1 - .04 \times 7)$
 - ▶ in state $(1,2)$, two total rewards 0.76 and 0.84
 - ▶ in state $(1,3)$ two total rewards 0.80 and 0.88
-
- ▶ Estimated utility: **mean** of sampled values
 - ▶ $U(1,1) = 0.72, U(1,2) = 0.80, U(1,3) = 0.84$ etc.

Direct utility estimation

- ▶ Assume utilities are independent (false).
It doesn't take into account that the utility of a state depends on the utilities of next states (the constraints given by the Bellman equations)

$$U^\pi(s) = R(s) + \gamma \sum_{s'} P(s'|s, \pi(s)) U^\pi(s')$$

- ▶ search in a much larger space
- ▶ convergence is very slow
- ▶ Have all episodes before

Content

Introduction

Passive learning

Model-based learning: Adaptive dynamic programming

Model-free learning: Direct utility estimation

Model-free learning: Temporal-Difference learning

Active learning

Q-learning

Deep Q-learning

Conclusions

Temporal-Difference learning

- ▶ Combines the advantages of *Adaptive Dynamic Programming* and *Direct utility estimation*
 - ▶ approximately satisfy the Bellman equations
 - ▶ update only the directly affected states
- ▶ Goal: estimate utilities $U^\pi(s)$, given the episodes generated using policy π . Actions are decided by policy π .
- ▶ The utilities are adjusted after each observed transition.

Example:

- ▶ After 1st trial: the estimates $U^\pi(1, 3) = 0.84$, $U^\pi(2, 3) = 0.92$.
- ▶ 2nd trial: the transition $(1, 3) \rightarrow (2, 3)$.
The constraint given by the Bellman equation requires the update of $U^\pi(1, 3)$.

Temporal-Difference learning

- ▶ The **temporal difference equation** uses the difference of utilities between successive states:

$$U^\pi(s) \leftarrow U^\pi(s) + \alpha(R(s) + \gamma U^\pi(s') - U^\pi(s))$$

α learning rate

The update involves only the successor s' , while the equilibrium conditions (Bellman eq.) involve all possible next states.

- ▶ The method applies a series of corrections to converge
- ▶ Obs: the method doesn't need a transitions model to perform the updates

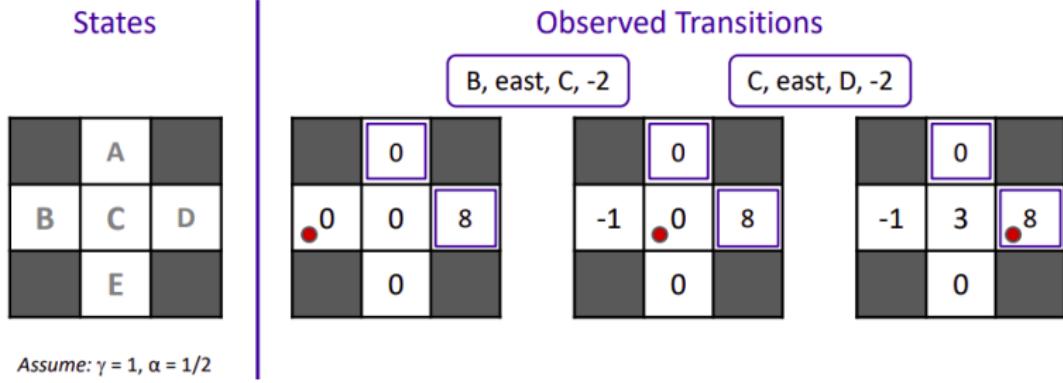
Temporal difference learning: pseudocode

```
function PASSIVE-TD-AGENT(percept) returns an action
    inputs: percept, a percept indicating the current state  $s'$  and reward signal  $r'$ 
    persistent:  $\pi$ , a fixed policy
         $U$ , a table of utilities, initially empty
         $N_s$ , a table of frequencies for states, initially zero
         $s, a, r$ , the previous state, action, and reward, initially null

    if  $s'$  is new then  $U[s'] \leftarrow r'$ 
    if  $s$  is not null then
        increment  $N_s[s]$ 
         $U[s] \leftarrow U[s] + \alpha(N_s[s])(r + \gamma U[s'] - U[s])$ 
    if  $s'.\text{TERMINAL?}$  then  $s, a, r \leftarrow \text{null}$  else  $s, a, r \leftarrow s', \pi[s'], r'$ 
    return  $a$ 
```

Demo: https://cs.stanford.edu/people/karpathy/reinforcejs/gridworld_td.html

Temporal difference learning: example

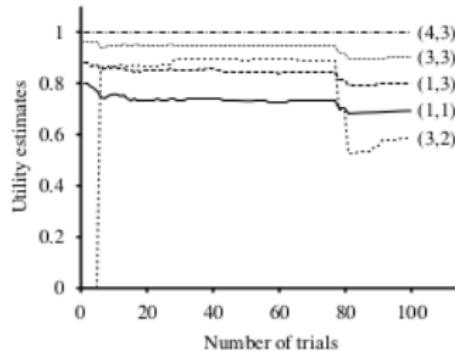
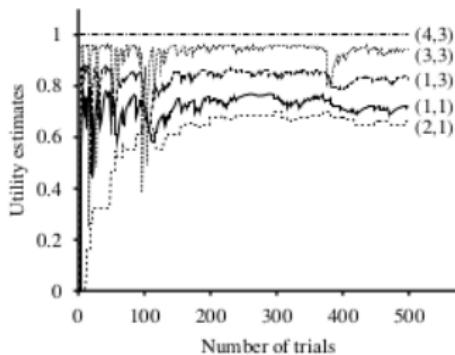


Temporal-Difference learning

- ▶ The learning rate α determines the speed of convergence to the true utility
- ▶ The mean value of $U^\pi(s)$ will converge to the correct value
 - ▶ many trials, sparse transitions occur rarely
 - ▶ if α is a function which decreases as the no. of visits to a state increases, then $U^\pi(s)$ converges to the correct value
 - ▶ $\alpha(n) = 1/n$ or $\alpha(n) = 1/(1 + n) \in (0, 1]$

Temporal differences learning vs. Adaptive dynamic programming

- ▶ TD needs no model, ADP is model-based
- ▶ TD use only the observed successor for updating and not all successors
- ▶ TD converges more slowly but performs simpler calculations



- ▶ TD can be seen as an approximation of ADP

Content

Introduction

Passive learning

Model-based learning: Adaptive dynamic programming

Model-free learning: Direct utility estimation

Model-free learning: Temporal-Difference learning

Active learning

Q-learning

Deep Q-learning

Conclusions

Passive learning vs. Active learning

- ▶ Passive agent has a fixed policy vs.
the active agent must **decide actions**
- ▶ Passive agent learns (transition probabilities and) utilities and chooses optimal actions

vs.

The active agent **updates its policy** as it learns

- ▶ the goal is to learn the optimal policy
- ▶ but, the utility function is known approximately

Exploitation vs. exploration

Exploitation-exploration dilemma of the agent

- ▶ to maximize his utility, based on the current knowledge, or
- ▶ to improve his knowledge

A compromise between

- ▶ exploitation
 - ▶ the agent stops learning and executes the actions given by the policy
 - ▶ exploration
 - ▶ the agent learns by trying new actions
- is necessary.

Exploitation-exploration dilemma: solutions

ϵ -greedy method

- ▶ Let $\epsilon \in [0, 1]$
- ▶ The next action to be selected will be:
 - ▶ a random action, with probability ϵ
 - ▶ the optimal action, with probability $1 - \epsilon$
- ▶ Implementation
 - ▶ initially, $\epsilon = 1$ (exploration)
 - ▶ when a learning episode ends, ϵ decreases (for ex. with 0.05) - the exploitation rate progressively increases
 - ▶ ϵ never falls below a threshold, e.g. 0.1
 - ▶ the agent always has a chance to explore, to avoid local optima

Content

Introduction

Passive learning

Model-based learning: Adaptive dynamic programming

Model-free learning: Direct utility estimation

Model-free learning: Temporal-Difference learning

Active learning

Q-learning

Deep Q-learning

Conclusions

Algorithm Q-Learning (Watkins, 1989)

- ▶ The Q-Learning algorithm learns an action-value function $Q(s, a)$ (*Q quality*).

$Q(s, a)$ the value of using action a in state s .

The relationship between utilities and Q values: $U(s) = \max_a Q(s, a)$

- ▶ The true equations at equilibrium when the Q values are correct

$$Q(s, a) = R(s) + \gamma \sum_{s'} P(s'|s, a) \max_{a'} Q(s', a')$$

These can be used in an iterative process that computes the exact Q values.

Q-Learning algorithm

- ▶ A TD agent which learns a Q function does not need a probabilistic model $P(s'|s, a)$ (*model-free learning*).
- ▶ For each sample (s, a, s', r) , update the Q value.
The update equation for *TD Q-Learning*:

$$Q(s, a) = Q(s, a) + \alpha(R(s) + \gamma \max_{a'} Q(s', a') - Q(s, a))$$

(by executing action a in state s we reach the state s')

The learning coefficient α determines the speed of updating the estimates; usually, $\alpha \in (0, 1)$

Q-Learning algorithm: pseudocode

```
function Q-LEARNING-AGENT(percept) returns an action
    inputs: percept, a percept indicating the current state  $s'$  and reward signal  $r'$ 
    persistent:  $Q$ , a table of action values indexed by state and action, initially zero
         $N_{sa}$ , a table of frequencies for state-action pairs, initially zero
         $s, a, r$ , the previous state, action, and reward, initially null

    if TERMINAL?( $s$ ) then  $Q[s, \text{None}] \leftarrow r'$ 
    if  $s$  is not null then
        increment  $N_{sa}[s, a]$ 
         $Q[s, a] \leftarrow Q[s, a] + \alpha(N_{sa}[s, a])(r + \gamma \max_{a'} Q[s', a'] - Q[s, a])$ 
         $s, a, r \leftarrow s', \text{argmax}_{a'} f(Q[s', a'], N_{sa}[s', a']), r'$ 
    return  $a$ 
```

f exploration function

- ▶ Q-learning converges to an optimal policy
- ▶ Q-Learning is slower than ADP

Q-learning

Example 1

Pacman is in an unknown MDP where there are three states [A, B, C] and two actions [Stop, Go]. We are given the following samples generated from taking actions in the unknown MDP. For the following problems, assume $\gamma = 1$ and $\alpha = 0.5$.

- (a) We run Q-learning on the following samples:

s	a	s'	r
A	Go	B	2
C	Stop	A	0
B	Stop	A	-2
B	Go	C	-6
C	Go	A	2
A	Go	A	-2

What are the estimates for the following Q-values as obtained by Q-learning? All Q-values are initialized to 0.

$$Q(C, \text{Stop}) = ?, Q(C, \text{Go}) = ?$$

Example 2

<https://huggingface.co/learn/deep-rl-course/unit2/q-learning-example>

SARSA

- ▶ The update equation:

$$Q(s, a) = Q(s, a) + \alpha(R(s) + \gamma Q(s', a') - Q(s, a))$$

(s', a') the pair (next state, next action)

SARSA uses the TD approach: update the Q table after each step until the solution converges / max. no. of iterations.

- ▶ Example: Windy Gridworld

<http://www.incompleteideas.net/book/ebook/node64.html>

Content

Introduction

Passive learning

Model-based learning: Adaptive dynamic programming

Model-free learning: Direct utility estimation

Model-free learning: Temporal-Difference learning

Active learning

Q-learning

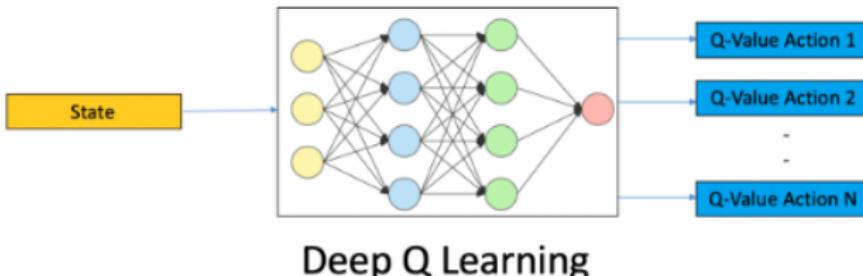
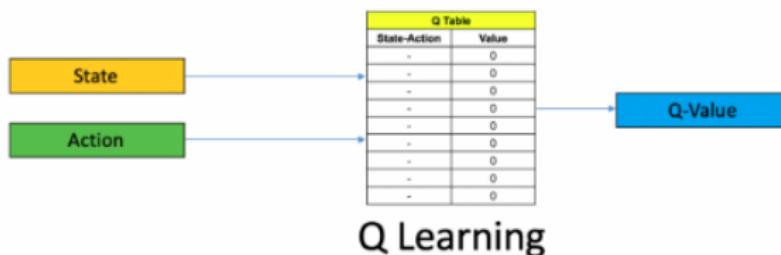
Deep Q-learning

Conclusions

Deep Reinforcement Learning

It uses a (deep) neural network to approximate the Q values

- ▶ input: a state
- output: an estimate of Q , for each possible action



Deep Reinforcement Learning

- ▶ Consider the update equation of the Q values:

$$\begin{aligned} Q(s, a) &= Q(s, a) + \alpha[r + \gamma \max_{a'} Q(s', a') - Q(s, a)] \\ &= (1 - \alpha)Q(s, a) + \alpha[r + \gamma \max_{a'} Q(s', a')] \end{aligned}$$

- ▶ The *loss function*: the mean squared error of the predicted Q value and the target value Q^* (not known).

Minimize $\text{loss}(s, a, s') = (Q(s, a) - \text{target}(s'))^2$, where
the target value: $\text{target}(s') = r + \gamma \max_{a'} Q(s', a')$

- ▶ Use *Gradient descent* to optimize the loss function.

Description: <https://deeplearningmath.org/deep-reinforcement-learning.html>
Demo: <https://cs.stanford.edu/people/karpathy/reinforcejs/>

Deep Q-learning

Problems:

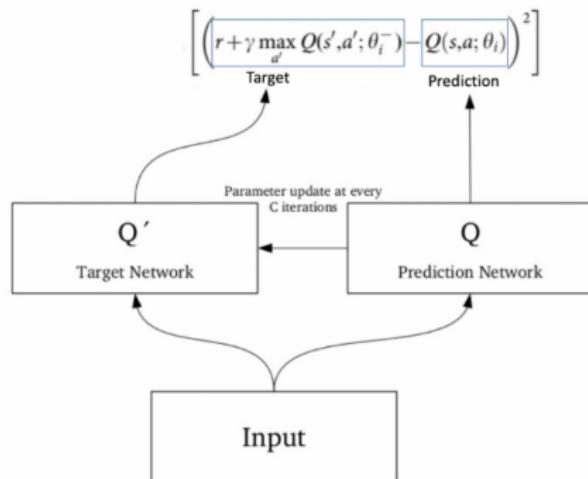
- ▶ the samples are correlated → the network cannot generalize
- ▶ $\text{target}(s')$ is an estimate → slow convergence/unstable alg.

Solutions:

- ▶ ϵ -greedy policy
- ▶ *experience replay*: store the experiences (s, a, r, s') and use them for training (*mini-batch*)

Double Deep Q-network

- ▶ the target value changes at each iteration; solution: a separate network to estimate the target value



- ▶ every C iterations, the parameters from the prediction network are copied to the target network

Function approximation

$$\hat{U}_\theta(s) = \theta_1 f_1(s) + \theta_2 f_2(s) + \dots \theta_n f_n(s)$$

f_1, \dots, f_n features

RL learn values for the parameters $\theta = \theta_1, \dots, \theta_n$ s.t the evaluation function \hat{U}_θ approximates the true utility function.

- ▶ updates the parameters after each trial
- ▶ use an error function and compute the gradients

$$E_j(s) = (\hat{U}_\theta(s) - u_j(s))^2 / 2$$

$u_j(s)$ the observed total reward from state s in trial j

$$\theta_i \leftarrow \theta_i - \alpha \frac{\delta E_j(s)}{\delta \theta_i} = \theta_i + \alpha(u_j(s) - \hat{U}_\theta(s)) \frac{\delta \hat{U}_\theta(s)}{\delta \theta_i}$$

- ▶ good generalization (visited states → unvisited states)

Content

Introduction

Passive learning

Model-based learning: Adaptive dynamic programming

Model-free learning: Direct utility estimation

Model-free learning: Temporal-Difference learning

Active learning

Q-learning

Deep Q-learning

Conclusions

Conclusions

- ▶ Reinforcement learning is necessary for agents that evolve in unfamiliar environments
- ▶ **Passive** learning evaluates a given policy
- ▶ **Active** learning learns an optimal policy

Bibliography

- ▶ *Artificial Intelligence: A modern Approach.* Ch. 21. Reinforcement Learning
- ▶ Sutton&Barto. *Reinforcement Learning.* An introduction
<http://incompleteideas.net/book/RLbook2020.pdf>